




Principle: route each subtask to the cheapest reliable executor; train TRMs on gates that are neither stable scripts nor safe LLM-only decisions